



AI reinvention made real

Accenture Innovation Challenge 2026

Team details

TEAM NAME:Dominator



Upendra Singh

College:IIT Kharagpur
Stream:Mechanical Engineering
Year of graduation:2028



Krishnakant Sahu

College:IIT Kharagpur
Stream:Computer Science & Engineering
Year of graduation:2027



Aditya Singh

College:IIT Kharagpur
Stream:Mechanical Engineering
Year of graduation:2028



CONTROLPLANE.AI – EXECUTIVE PROPOSITION

AI needs a control plane, not another guardrail.




ControlPlane.ai
The Missing Safety Layer for AI

ControlPlane.ai is the first evidence-backed control plane that continuously observes AI responses, validates risk evidence, issues time-bounded warrants, and turns oversight into policy decisions in real time — **without slowing AI down.**

WHAT WE DELIVER

- ✓ Proactive risk control across **Performance, Cost, and Responsibility** before issues reach the user.
- ✓ Evidence-backed warrants with expiry, refusal, and drift detection for governance you can **trust.**
- ✓ **Measurable impact** through higher risk detection, fewer incidents, lower rework, and faster, safer AI experiences.



BUSINESS VALUE AT A GLANCE



MORE RELIABLE

Catch issues before users act. Reduce wrong, harmful and unsafe outputs.



COST EFFICIENT

Spend expensive checks only where risk is high. Lower rework and operational waste.



FASTER EXPERIENCE

Keep latency low. Deep checks happen only when needed.



TRUST & COMPLIANCE

Built-in oversight for responsibility, safety and data privacy. Ready for regulators.



MEASURABLE IMPACT

Evidence-backed warrants, drift detection and full auditability you can prove.

EVIDENCE SO FAR (MEASURED)

0.826

TriviaQA Probe AUROC

0.813

Long-context Probe AUROC

0.794

Hinglish PII Recall

13

Warrant Cells VALID

4

Warrant Cells REFUSED

39

Warrant Cells UNVALIDATED

42.2%

Theoretical Abstention Floor (5% Residual Risk)




ControlPlane.ai turns AI oversight from after-the-fact discovery into real-time control you can **watch, catch, and act on** — at scale.

> 02 — PROBLEM & MARKET NEED


AI failures are rising. Traditional controls are stale, siloed, and discover problems too late.

Enterprise AI is moving faster than oversight.

Models hallucinate, leak data, produce biased or harmful content, and waste compute — and organizations usually find out only **after a user has already acted on it.**



THE PROBLEM IS REAL, WIDESPREAD, AND GROWING



Sources: 1. IBM 2. Gartner 3. Stanford AI Index 4. Deloitte 5. IBM Cost of a Data Breach 2024 6. Accenture 7. Deloitte Regulatory Outlook 2024 8. Grand View Research




THE GOVERNANCE GAP

Most organizations rely on siloed, after-the-fact controls — logs, audits, human reviews, and point solutions.

They tell you what happened.

ControlPlane.ai tells you what to do — in real time.

WHY EXISTING CONTROLS FALL SHORT



The market needs a real-time, evidence-backed control plane that continuously observes, validates, and decides — **before risks become incidents.**



ControlPlane.ai closes this gap.
It's the **missing safety layer** for AI.



ControlPlane.ai | turns AI oversight from after-the-fact discovery into real-time control you can **watch, catch, and act on** — at scale.

> 03 — CONTROLPLANE.AI SOLUTION

Detector → Envelope → Warrant → Policy.

From raw signals to trusted, time-bounded decisions.

ControlPlane.ai sits on top of any model and turns AI oversight into an **evidence-backed decision pipeline** that is continuous, auditable, and actionable.



We don't just **detect** risk.
We **validate** it, stand behind it with a **warrant**, and **decide** what to do — in real time.

THE CONTROLPLANE DECISION PIPELINE

1. DETECTOR (What might be wrong?)



Lightweight models and probes scan every response for risk signals across **Performance**, **Cost**, and **Responsibility**.



Risk signals with scores



2. EVALUATION ENVELOPE (How good is this evidence?)



We define the scope: data slice, time, model, thresholds, and methods. The envelope makes evidence comparable and fair.



Context & criteria for validation



3. WARRANT (What can we stand behind?)




ControlPlane issues a time-bounded warrant (or refuses) based on validated evidence. Includes uncertainty, expiry, and drift status.



Warrant or Refusal (with expiry)



4. POLICY DECISION (What should we do?)



Policy engine converts warrants into actions — Allow, Edit/Review, Block, or Escalate — with full auditability.



Action + Rationale + Audit log

Continuous feedback loop: outcomes feed re-validation and drift detection

WHAT MAKES CONTROLPLANE DIFFERENT



EVIDENCE FIRST

Every decision is backed by measured evidence, not assumptions.



TIME-BOUNDED

Warrants have expiry. Riskiest assumption is treated as data.



REFUSAL IS A FEATURE

We refuse when evidence is insufficient. Silence is safer than false trust.



DRIFT-AWARE

Continuously monitors and re-validates as models and data evolve.



ACTIONABLE & AUDITABLE

Every action is explainable, traceable, and ready for compliance.

OUTCOMES YOU CAN EXPECT



30–60%

fewer high-risk incidents



20–40%

lower compute and rework costs



2–5x

faster risk review with automation



90%+

policy decisions with full audit trail



70%+

stakeholder trust in AI outcomes



ControlPlane.ai brings together detection, validation, and decisioning in a single control plane so organizations can **trust AI** — and **act with confidence**.



Works with any model and any use case



Model-agnostic, cloud-agnostic



Built for scale, built for enterprise



ControlPlane.ai

turns AI oversight from after-the-fact discovery into real-time control you can **watch, catch, and act on** — at scale.

> 04 — HOW IT WORKS TECHNICALLY

Four-tier architecture: maximum coverage at minimal cost and latency.

ControlPlane.ai uses a progressive **four-tier pipeline** to observe every AI response, trigger deeper checks only when needed, and make policy decisions in real time.



Cheap signals first.
Expensive verdicts only when the probe says it matters.

KEY DESIGN PRINCIPLES



1. **Trigger, not verdict**
Tiers 0/1 only decide where to spend a judge call.




2. **One probe, two jobs**
Monitor all traffic and route hard questions before gen.



3. **Open weights unlock Tier 2**
Tiers 0/1/3 can run on closed APIs. Tier 2 must be open and hostable.

THE FOUR-TIER ARCHITECTURE



ARCHITECTURE FLOW

100% Traffic → Cheap Signals →
Small Uncertain Slice → Deep Verdict



1. **Observe & Trigger**
All responses are observed. Probe triggers deeper checks only for risky cases.



2. **Route & Adjudicate**
Hard questions are routed to Tier 2 and then Tier 3 before reaching the user.



3. **Decide & Act**
Policy decisions are executed in real time with full audit and rationale.

TYPICAL TRAFFIC DISTRIBUTION



● Tiers 0-2 (automated) 95-97%
● Tier 3 (deep adjudication) 3-5%

LATENCY SNAPSHOT (P50)



OPTIMIZED FOR



Low latency



High coverage



Cost efficiency



Strong governance

MATCHED-BUDGET COMPARISON

Policy	Judge Calls	Responses Seen by Any Check	Errors Caught	Relative Cost
Judge everything	1.00N	100%	~100%	20x
Random sample 5%	0.05N	5%	~5%	1x
Probe-triggered (ControlPlane)	~0.05N	100%	(YOUR MEASURED R)	1x

WHY THIS ARCHITECTURE WORKS



MAX COVERAGE

Every response is seen by cheap checks; hard cases get deep review.



MIN COST

Expensive adjudication used only where the probe says it matters.



LOW LATENCY

Tiered design keeps most traffic fast and non-blocking.



FULL GOVERNANCE

Decisions are explainable, auditable, and evidence-backed.



MODEL AGNOSTIC

Works with any model, any provider, any deployment.



> 05 — WHY IT IS DIFFERENT

Evidence + expiry + refusal + drift.
That's why ControlPlane.ai decisions are trusted.

Most solutions detect risk or log incidents. ControlPlane.ai goes further — it **validates** evidence, issues time-bound **warrants**, knows when to **refuse**, and **continuously adapts to drift**.



We don't just raise alerts. We provide **trusted evidence** with context, validity, and expiry.

CONTROLPLANE.AI vs. TRADITIONAL APPROACHES

	 EVIDENCE, NOT ASSUMPTION	 TIME-BOUND WARRANTS	 REFUSAL IS A FEATURE	 DRIFT-AWARE BY DESIGN	 ACTIONABLE & AUDITABLE
	Decisions backed by measured evidence, not guesses.	Every warrant has an expiry based on data freshness & stability.	If evidence is insufficient or conflicting, we refuse to decide.	Continuously monitors performance & data drift in production.	Every decision is explainable, traceable, and audit-ready.
CONTROLPLANE.AI	Evidence-backed with confidence & context.	Warrants with expiry & validity window.	Warrant REFUSED when evidence is insufficient.	Detects drift and re-validates warrants automatically.	Full audit trail: evidence → warrant → policy → action.
TRADITIONAL APPROACHES	Heuristic or rule-based assumptions.	Static models, no concept of validity window.	Always returns a score or decision, even when unsure.	Drift detected late (after incidents) or not at all.	Limited visibility; hard to audit or explain decisions.

BUSINESS IMPACT OF THE DIFFERENCE



FEWER SURPRISES

Catch issues before users act, preventing bad outcomes.



LOWER COST & REWORK

Spend compute and reviewer time only where it matters.



FASTER, SAFER EXPERIENCES

Low-latency decisions with high confidence and coverage.



STRONGER GOVERNANCE

Built-in traceability, accountability, and regulatory readiness.



COMPOUNDING VALUE

Better data & feedback → stronger warrants → smarter decisions.

THE CONTROLPLANE ADVANTAGE



OBSERVE

Capture signals in real time.



VALIDATE

Validate across critical dimensions.



WARRANT

Issue or refuse with expiry.



DECIDE

Convert warrant into policy.



ACT & LEARN

Execute, collect feedback and re-validate.

PROOF POINTS (MEASURED)

0.826

TriviaQA Probe AUROC
[0.793–0.857]

0.813

Long-context Probe AUROC
[0.780–0.845]

0.794

Hinglish PII Reference Recall
[0.698–0.885]

13

Warrant Cells VALID

4

Warrant Cells REFUSED

39

Warrant Cells UNVALIDATED

42.2%

Theoretical Abstention Floor (5% Residual Risk)



> 06 — TARGET CUSTOMERS

From Customer Support → Internal Knowledge → Decision Support → Governance.
One control plane. Many high-impact use cases.



Our mission: Protect every high-stakes AI interaction across the enterprise.
ControlPlane.ai **adapts to the context, risk, and needs** of every function.

1. CUSTOMER SUPPORT AI for front-line agents



- Agent copilot, auto-replies, summarization, intent classification
- Handles PII, refunds, complaints, and sensitive interactions

KEY OUTCOMES

- ✓ Fewer harmful / incorrect responses
- ✓ Lower escalations & rework
- ✓ Better CSAT and agent productivity

2. INTERNAL KNOWLEDGE AI for employees



- Knowledge search, Q&A, document summarization
- Enterprise data, policies, and proprietary content

KEY OUTCOMES

- ✓ Accurate, trustworthy answers
- ✓ Prevent data leakage & policy violations
- ✓ Higher employee productivity

3. DECISION SUPPORT AI for business decisions



- Analysis, recommendations, what-if scenarios
- Impacts revenue, risk, compliance, and strategic outcomes

KEY OUTCOMES

- ✓ Higher decision quality & consistency
- ✓ Reduced analytical risk
- ✓ Stronger auditability of decisions

4. GOVERNANCE & COMPLIANCE AI across the enterprise



- AI inventory, risk oversight, continuous monitoring
- Regulatory compliance, model risk management

KEY OUTCOMES

- ✓ Evidence for regulators & auditors
- ✓ Standardized risk controls
- ✓ Enterprise-wide AI governance

WHO BUYS & CHAMPIONS



Business Leaders (Buyers)

COO, CRO, CDO, Head of Operations, Business Unit Heads
Care about outcomes, risk, efficiency, and customer trust.



Risk & Compliance (Champions)

CRO, CISO, Head of Risk, Compliance Officers
Ensure safe, responsible, and auditable AI deployments.



Technology Leaders

CTO, Head of AI/ML, Engineering Leads
Need scalable, low-latency, easy-to-integrate controls.



Operations & Support Leaders

Head of Customer Support, Knowledge Ops, Service Ops
Drive adoption, measure impact, and improve day-to-day outcomes.

KEY ENTERPRISE USE CASES (EXAMPLES)



Chat / Copilot

Safe responses, no hallucinations or policy violations



PII & Secret Protection

Detect and block or mask sensitive data in real time



Policy & Compliance

Enforce policies, prevent prompt injection & jailbreaks



Quality & Fact Check

Validate claims, reduce confidently wrong answers



Cost & Efficiency

Catch waste, redundant processing, and unnecessary rework



Drift & Change Detection

Monitor performance and data drift continuously

PROOF POINTS (CURRENT RESULTS)

0.826

TriviaQA Probe
AUROC
[0.793–0.857]

0.813

Long-context Probe
AUROC
[0.780–0.845]

0.794

Hinglish PII
Reference Recall
[0.698–0.885]

13

Warrant Cells
VALID

4

Warrant Cells
REFUSED

39

Warrant Cells
UNVALIDATED

42.2%

Theoretical Abstention
Floor (5% Residual Risk)



> 07 — BUSINESS CASE

Risk reduction + Review efficiency + Governance + Low latency.
ControlPlane.ai delivers measurable value across the enterprise.



AI can be confidently wrong, biased, costly, or non-compliant.
ControlPlane.ai prevents issues before users act — reducing risk, improving efficiency, and ensuring trusted, governed AI at scale.



30–60%
fewer high-risk incidents



2–5x
faster risk review with automation



20–40%
lower compute & rework costs



90%+
policy decisions with full audit trail

01

02

03

04

05

06

07

08

09

10

11

12

13

14

15

16

17

18

19

20

21

22

23

24

25

26

27

28

29

30

31

32

33

34

35

36

37

38

39

40

41

42

43

44

45

46

47

48

49

50

51

52

53

54

55

56

57

58

59

60

61

62

63

64

65

66

67

68

69

70

71

72

73

74

75

76

77

78

79

80

81

82

83

84

85

86

87

88

89

90

91

92

93

94

95

96

97

98

99

100

101

102

103

104

105

106

107

108

109

110

111

112

113

114

115

116

117

118

119

120

121

122

123

124

125

126

127

128

129

130

131

132

133

134

135

136

137

138

139

140

141

142

143

144

145

146

147

148

149

150

151

152

153

154

155

156

157

158

159

160

161

162

163

164

165

166

167

168

169

170

171

172

173

174

175

176

177

178

179

180

181

182

183

184

185

186

187

188

189

190

191

192

193

194

195

196

197

198

199

200

201

202

203

204

205

206

207

208

209

210

211

212

213

214

215

216

217

218

219

220

221

222

223

224

225

226

227

228

229

230

231

232

233

234

235

236

237

238

239

240

241

242

243

244

245

246

247

248

249

250

251

252

253

254

255

256

257

258

259

260

261

262

263

264

265

266

267

268

269

270

271

272

273

274

275

276

277

278

279

280

281

282

283

284

285

286

287

288

289

290

291

292

293

294

295

296

297

298

299

300

301

302

303

304

305

306

307

308

309

310

311

312

313

> 08 — PILOT PLAN

One workload → shadow mode → controlled actioning.
Prove value fast. Scale with confidence.



Goal: Validate ControlPlane.ai on a real production workload, measure outcomes, and progressively take action with guardrails — **without disrupting the business.**



Validate
in production



Quantify risk
reduction



Improve efficiency
& latency



Build governance
that scales

PILOT JOURNEY: FOUR PHASES

1. SELECT WORKLOAD (Week 0–1)



- Pick one high-impact, bounded use case (e.g., Customer Support Copilot).
- Define success metrics and risk appetite.
- Integrate via LiteLLM proxy in observation mode.

DELIVERABLES

- ✓ Use case & scope document
- ✓ Risk register & policy draft
- ✓ Instrumentation & dashboards

2. SHADOW MODE (OBSERVE) (Week 2–4)



- ControlPlane observes 100% of traffic.
- No user impact. No blocking.
- Collect data, issue warrants, and measure drift & uncertainty.

DELIVERABLES

- ✓ Baseline risk & warrant report
- ✓ Drift & uncertainty analysis
- ✓ Cost, latency & flag rate report

3. CONTROLLED ACTIONING (LIMITED) (Week 5–8)



- Take automated action on low-risk, high-confidence cases.
- Escalate uncertain/high-risk to human review.
- Tune policies and thresholds.

DELIVERABLES

- ✓ Action decision outcomes
- ✓ Review workload reduction report
- ✓ Accuracy & false positive analysis

4. EXPAND & OPTIMIZE (Week 9–12+)



- Expand to more intents, channels, and adjacent workflows.
- Optimize cost, latency, and policies.
- Operationalize monitoring & governance.

DELIVERABLES

- ✓ Full pilot report & ROI estimate
- ✓ Policy library & runbook
- ✓ Scale-up plan & roadmap

PILOT PRINCIPLES



SAFE BY DESIGN

Shadow first. Action later. Always human-in-the-loop for uncertain or high-risk.



EVIDENCE-BACKED

Decisions backed by warrants with validity, expiry, and context.



TIME-BOUND & ADAPTIVE

Continuously re-validate as data and models drift.



TRANSPARENT & AUDITABLE

Full traceability: evidence → warrant → policy → action.



MEASURABLE IMPACT

Every phase tied to clear metrics and business value.

PILOT GUARDRAILS



Never block without high confidence & valid warrant



Always escalate when uncertain



PII / secrets never leave tenant



All actions are logged, explainable & reversible

SUCCESS METRICS (MEASURED DURING PILOT)



RISK REDUCTION

30–60%

reduction in high-risk incidents



REVIEW EFFICIENCY

2–5x

faster review with automation



COMPUTE & REWORK SAVINGS

20–40%

lower compute and rework costs



LATENCY IMPACT

P50 < 150ms

at 95–97% of traffic



GOVERNANCE COVERAGE

90%+

policies with full audit trail



ADOPTION READINESS

>70%

stakeholder trust in AI outcomes

PROOF POINTS (CURRENT RESULTS)

0.826

TriviaQA Probe AUROC
[0.793–0.857]

0.813

Long-context Probe AUROC
[0.780–0.845]

0.794

Hinglish PII
Reference Recall
[0.698–0.885]

13

Warrant Cells
VALID

4

Warrant Cells
REFUSED

39

Warrant Cells
UNVALIDATED

42.2%

Theoretical Abstention
Floor (5% Residual Risk)



> 09 — ROADMAP

Evidence hardening → gateway → actioning → enterprise → continuous certification.
A phased path to trusted, scalable AI oversight.



> 10 — RISKS & MITIGATIONS

We identify risks early, mitigate by design, and continuously adapt.
Safe AI outcomes — with evidence, control, and accountability.

RISK CATEGORY	DESCRIPTION	POTENTIAL IMPACT	LIKELIHOOD	MITIGATIONS	RESIDUAL RISK
 1. HALLUCINATIONS & UNSAFE OUTPUTS	Model generates incorrect, misleading, or harmful responses.	Wrong decisions, customer harm, compliance exposure, brand damage.	● MEDIUM	✓ Evidence-backed warrants with validity & expiry ✓ Refusal on insufficient/ conflicting evidence ✓ Guardrails & policy checks pre-action	● LOW
 2. DATA & PRIVACY EXPOSURE	PII, secrets, or sensitive data leak through prompts, outputs, or logs.	Regulatory fines, legal action, loss of customer trust.	● MEDIUM	✓ PII/secret detection & redaction ✓ Data minimization & encryption ✓ Strict access controls & tenancy isolation	● LOW
 3. POLICY BYPASS & JAILBREAKS	Malicious prompts or techniques attempt to bypass controls.	Policy violations, unsafe actions, compliance breach.	● HIGH	✓ Policy engine with allow/edit/block & escalation ✓ Jailbreak detection & response ✓ Continuous red-teaming & evals	● MEDIUM
 4. MODEL DRIFT & DEGRADATION	Model behavior changes over time, reducing reliability.	Lower accuracy, more refusals, inconsistent outcomes.	● MEDIUM	✓ Continuous monitoring & drift detection ✓ Re-validation of warrants & datasets ✓ Auto recalibration & model reviews	● LOW
 5. INTEGRATION & OPERATIONAL RISK	Failures in integrations, configs, or dependencies impact availability.	Service disruption, latency spikes, failed actions.	● MEDIUM	✓ Resilient architecture & fallbacks ✓ Health checks, alerts & runbooks ✓ Change control & staged rollouts	● LOW
 6. HUMAN & PROCESS RISK	Misuse, misconfiguration, or unclear ownership and processes.	Inconsistent outcomes, shadow AI, audit findings.	● MEDIUM	✓ RBAC, least privilege & approvals ✓ Training, playbooks & clear ownership ✓ Audit trails & accountability by design	● LOW

OVERALL RISK POSTURE

 **LOW-MEDIUM**

With designed controls, continuous monitoring, and governance.

RISK MANAGEMENT PRINCIPLES

 **PREVENT**
Stop issues before users act.

 **DETECT**
Continuously monitor and surface risks.

 **RESPOND**
Escalate, remediate, and learn fast.

 **IMPROVE**
Strengthen controls through evidence and continuous feedback.



PROOF POINTS (CURRENT RESULTS)	0.826 TriviaQA Probe AUROC [0.793–0.857]	0.813 Long-context Probe AUROC [0.780–0.845]	0.794 Hinglish PII Reference Recall [0.698–0.885]	13 Warrant Cells VALID	4 Warrant Cells REFUSED	39 Warrant Cells UNVALIDATED	42.2% Theoretical Abstention Floor (5% Residual Risk)
-----------------------------------	--	--	--	------------------------------	-------------------------------	------------------------------------	---

> 11 — COMMERCIAL MODEL

Outcome-aligned, usage-aware pricing.
Start small. Scale predictably. Realize value fast.



We align price to value: protect more interactions, reduce risk, and improve efficiency.
Flexible options for teams of any size — with transparent, predictable costs.



Outcome-aligned
Pay for value delivered



Usage-aware
Scale with confidence



Predictable
Simple & transparent



Enterprise-ready
Secure & compliant

SUBSCRIPTION EDITIONS

	STARTER Pilot & small teams	PROFESSIONAL Growing teams & use cases	BUSINESS Enterprise scale	ENTERPRISE Global & highly regulated
BEST FOR	Pilot projects, single workload, small teams	Multiple use cases, team collaboration, controlled actioning	Enterprise-wide deployments, integrations, SSO	Global scale, multi-region, regulated industries
INCLUDES	Up to 50K interactions/month 1 workload Core controls Standard dashboards	Up to 500K interactions/month Up to 5 workloads All core capabilities Advanced dashboards	Up to 5M interactions/month Unlimited workloads All capabilities Role-based access	Custom interactions Unlimited everything Dedicated environment Advanced compliance
KEY CAPABILITIES	✓ Detector & Envelope ✓ Warrant (evidence + expiry) ✓ Policy (allow / block / warn) ✓ Shadow mode ✓ Basic reporting	✓ Everything in Starter ✓ Controlled actioning ✓ Drift detection ✓ Advanced analytics ✓ Priority alerts & notifications	✓ Everything in Professional ✓ Enterprise integrations (SIEM, ITSM) ✓ Policy engine (ABAC) ✓ Governance workspace ✓ SLA & advanced reporting	✓ Everything in Business ✓ Custom detectors & policies ✓ Data residency & isolation ✓ Advanced audit & attestations ✓ Customer success & TAM
SUPPORT	Email support Business hours	Email + chat support Business hours	24x7 support Response SLA	24x7 premium support Dedicated TAM
PRICING (Annual)	\$10K / year Up to 50K interactions/month Overage: \$0.20 / 1K interactions	\$50K / year Up to 500K interactions/month Overage: \$0.15 / 1K interactions	\$150K / year Up to 5M interactions/month Overage: \$0.10 / 1K interactions	Custom Tailored to your needs Volume & commitments

USAGE IS SIMPLE & TRANSPARENT



We bill on monthly interactions processed by ControlPlane.ai.

One interaction = one AI response evaluated end-to-end.

ADD-ONS (OPTIONAL)



Data Retention Packs
Extended retention for evidence beyond included period



Integrations Pack
Additional connectors, custom integrations & professional services



Managed Services
Policy as a service, 24x7 monitoring, & incident response

PRICING PRINCIPLES

- ✓ Start small, prove value, scale with confidence
- ✓ No surprises. No hidden fees.
- ✓ Aligns to outcomes that matter: risk reduction, efficiency, trust

COMMERCIAL GUARANTEES



30-day success checkpoint
Ensure value & alignment or course-correct



Pilot guarantee
Results or money back for eligible pilots



Price protection
No price increases for 12 months



Scale discounts
Volume tiers for predictable unit economics



Cancel anytime
No long-term lock-in for subscriptions

PROOF POINTS (CURRENT RESULTS)

0.826
TriviaQA Probe AUROC
[0.793–0.857]

0.813
Long-context Probe AUROC
[0.780–0.845]

0.794
Hinglish PII
Reference Recall
[0.698–0.885]

13
Warrant Cells
VALID

4
Warrant Cells
REFUSED

39
Warrant Cells
UNVALIDATED

42.2%
Theoretical Abstention
Floor (5% Residual Risk)

12 — ACCENTURE OPPORTUNITY / SCALE

Scale trusted AI oversight across clients, industries, and geographies.
Accelerate value for clients. New growth for Accenture.



Accenture + ControlPlane.ai: combine deep industry and transformation expertise with purpose-built AI safety and governance — to deliver trusted, auditable AI at scale.



Global scale
& delivery



Industry depth
& trust



AI governance
leadership



Cloud & data
at scale



Responsible AI
by design

1. LARGE & GROWING MARKET

AI oversight is now mission-critical



- Exploding AI adoption across every industry and function.
- Regulation and risk driving demand for governance and controls.
- Early innings — massive addressable opportunity.

ACCENTURE ADVANTAGE

- ✓ Global client base & industry reach
- ✓ Trusted advisor & transformation partner
- ✓ Scale to enterprise & multi-geo

2. STRATEGIC DIFFERENTIATOR

The missing safety layer for AI



- Complements Accenture's AI services and platforms.
- Embedded safety, auditability, and accountability.
- Strengthens Responsible AI leadership in the market.

ACCENTURE ADVANTAGE

- ✓ Differentiated, IP-led offering
- ✓ Elevates trust in AI outcomes
- ✓ Accelerates deal wins & retention

3. HIGH-VALUE CLIENT OUTCOMES

Measurable risk reduction & efficiency



- Fewer incidents, lower risk, better compliance.
- Faster reviews, lower cost, higher throughput.
- Stronger governance and stakeholder confidence.

ACCENTURE ADVANTAGE

- ✓ Outcome-based engagements
- ✓ Clear ROI & business case
- ✓ Expands advisory & managed services

4. SCALE THROUGH ACCENTURE

Delivery, assets, and ecosystem



- Leverage Accenture delivery scale, capabilities & methodologies.
- Pre-built assets, connectors, and industry policies.
- Ecosystem partnerships & cloud marketplaces.

ACCENTURE ADVANTAGE

- ✓ Faster time-to-value at scale
- ✓ Reusable assets & IP leverage
- ✓ Ecosystem reach & co-sell

5. RECURRING & EXPANDABLE

From projects to platform



- Land-and-expand across workloads, functions, and geographies.
- Drive platform adoption and data network effects.
- Long-term managed services & subscriptions.

ACCENTURE ADVANTAGE

- ✓ Recurring revenue & high retention
- ✓ Cross-sell & up-sell opportunities
- ✓ Platform-driven scale economics

WHERE WE PLAY



Industries

FSI, Health, Public Service, Retail, Manufacturing, Energy, Comms & more



Client Segments

Enterprise, Mid-market, Public sector, Global 2000



Geographies

Global scale with local delivery — powered by Accenture



Use Cases

Customer support, underwriting, claims, operations, code-gen, copilots, and more

PATH TO SCALE TOGETHER

- 1 Co-innovate & co-invest in roadmap & assets
- 2 Joint GTM & thought leadership
- 3 Industry solutions & accelerators
- 4 Global delivery & managed services
- 5 Marketplace listings & ecosystem scale

JOINT VALUE DRIVERS



30–60%

reduction in high-risk
AI incidents



2–5x

faster review with
automation



20–40%

lower compute &
rework costs



P50 < 150ms

latency at 95–97%
of traffic



90%+

policies with full
audit trail



>70%

stakeholder trust in
AI outcomes

PROOF POINTS (CURRENT RESULTS)

0.826

TriviaQA Probe AUROC
[0.793–0.857]

0.813

Long-context Probe AUROC
[0.780–0.845]

0.794

Hinglish PII
Reference Recall
[0.698–0.885]

13

Warrant Cells
VALID

4

Warrant Cells
REFUSED

39

Warrant Cells
UNVALIDATED

42.2%

Theoretical Abstention
Floor (5% Residual Risk)



13 — FINAL ASK

Partner with us to make AI safe, auditable, and accountable — at scale.
Let's build the missing safety layer for AI, together.



Our ask:

Partner with Accenture to take ControlPlane.ai to enterprise scale and global impact. Together, we can set the standard for safe, trustworthy, and accountable AI.



Reduce risk



Increase trust



Improve efficiency



Accelerate value



Lead the market

1

WHAT WE ARE ASKING FOR



- Strategic partnership to scale ControlPlane.ai with Accenture globally.
- Co-investment to accelerate product roadmap, integrations, and ecosystem.
- Go-to-market collaboration across industries, regions, and client segments.
- Joint IP creation and thought leadership to set the industry standard.

WHAT YOU GET

- ✓ Differentiated IP & proprietary capability
- ✓ New revenue streams & high retention
- ✓ Stronger client outcomes & trust
- ✓ Leadership position in Responsible AI

2

HOW WE WILL WORK TOGETHER



- Joint governance and operating model.
- Aligned investment & prioritized roadmap.
- Co-developed industry solutions and reference architectures.
- Shared GTM, delivery, and success plans.
- Continuous measurement and value realization.

OUR COMMITMENT

- ✓ Open, transparent, and collaborative
- ✓ Outcome-focused and measurable
- ✓ Built for security, compliance, and scale
- ✓ Always accountable for impact

3

WHAT WE NEED TO GET STARTED



- Executive sponsorship to align and commit.
- Nominate a joint pod (business, delivery, and tech leads).
- Greenlight pilot(s) in priority workload(s).
- Access to integration environments and data (as applicable).
- Co-investment to accelerate roadmap and scale.

IMMEDIATE NEXT STEPS

- ✓ Executive alignment & success criteria (Week 1)
- ✓ Pilot scoping & joint plan (Weeks 1–2)
- ✓ Shadow mode launch (Weeks 2–4)
- ✓ Review & scale decision (Weeks 6–8)

WHY NOW



Regulation and risk are accelerating.



AI adoption is scaling — trust is the bottleneck.



Clients demand evidence, control, and outcomes.



Early leadership creates durable advantage.

THE OPPORTUNITY TOGETHER



Define and own the AI safety category.



Deliver measurable outcomes across every Accenture capability.



Scale globally across industries and geographies.



Create long-term value for clients and Accenture.

FINAL WORD



ControlPlane.ai is the missing safety layer for AI. With Accenture, we can make every AI interaction safe, trusted, and accountable — at enterprise scale.
Let's build the future of responsible AI, together.

EXPECTED IMPACT TOGETHER



30–60%

reduction in high-risk AI incidents



2–5x

faster review with automation



20–40%

lower compute & rework costs



P50 < 150ms

latency at 95–97% of traffic



90%+

policies with full audit trail



>70%

stakeholder trust in AI outcomes

PROOF POINTS (CURRENT RESULTS)

0.826

TriviaQA Probe AUROC
[0.793–0.857]

0.813

Long-context Probe AUROC
[0.780–0.845]

0.794

Hinglish PII
Reference Recall
[0.698–0.885]

13

Warrant Cells
VALID

4

Warrant Cells
REFUSED

39

Warrant Cells
UNVALIDATED

42.2%

Theoretical Abstention
Floor (5% Residual Risk)

Thank you

