



AI reinvention made real

Accenture Innovation Challenge 2026

Team details

TEAM NAME:Dominator



Upendra Singh

College:IIT Kharagpur
Stream:Mechanical Engineering
Year of graduation:2028



Krishnakant Sahu

College:IIT Kharagpur
Stream:Computer Science & Engineering
Year of graduation:2027



Aditya Singh

College:IIT Kharagpur
Stream:Mechanical Engineering
Year of graduation:2028

> CONTROLPLANE.AI – THE MISSING SAFETY LAYER FOR AI

AI can be confidently wrong, quietly expensive, or subtly biased —
and organizations usually find out only after a user has already acted on it.



Every AI deployment carries the same risk: it can be confidently wrong, quietly expensive, or subtly biased – usually discovered only **after a user has already acted on it.**

ControlPlane Checker is a technology layer that sits on top of any model and turns AI oversight from an **after-the-fact** discovery into something you can **watch, catch, and act on live.**

AI FAILURES ARE EXPENSIVE

- Financial loss**
due to wrong decisions
- Reputational damage**
from unsafe or biased outputs
- Operational delays**
caused by rework and manual interventions

PERFORMANCE

Is it right or confidently wrong?

What can go wrong?

Hallucinations, incorrect answers, low quality outputs

COST

Is it burning more compute or rework than it should?

What can go wrong?

Excessive compute, redundant processing, unnecessary rework

RESPONSIBILITY

Is it biased, unsafe, or leaking data?

What can go wrong?

Bias, harmful content, privacy risks, data leakage

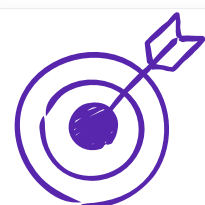
THE GOVERNANCE GAP

Today's AI controls are often fragmented and reactive.

- MONITOR:** dashboards show what happened — not always what is happening.
- DETECT:** risks are identified after an output reaches the user.
- DECIDE:** teams lack a consistent way to determine allow, edit, block, or escalate.
- ACT:** human intervention often happens only after business impact occurs.

ControlPlane.ai closes the gap by turning AI oversight into a real-time control loop.

OBSERVE → **DETECT** → **DECIDE** → **ACT**



The biggest AI risk isn't a wrong answer.
It's a wrong answer that nobody catches in time.

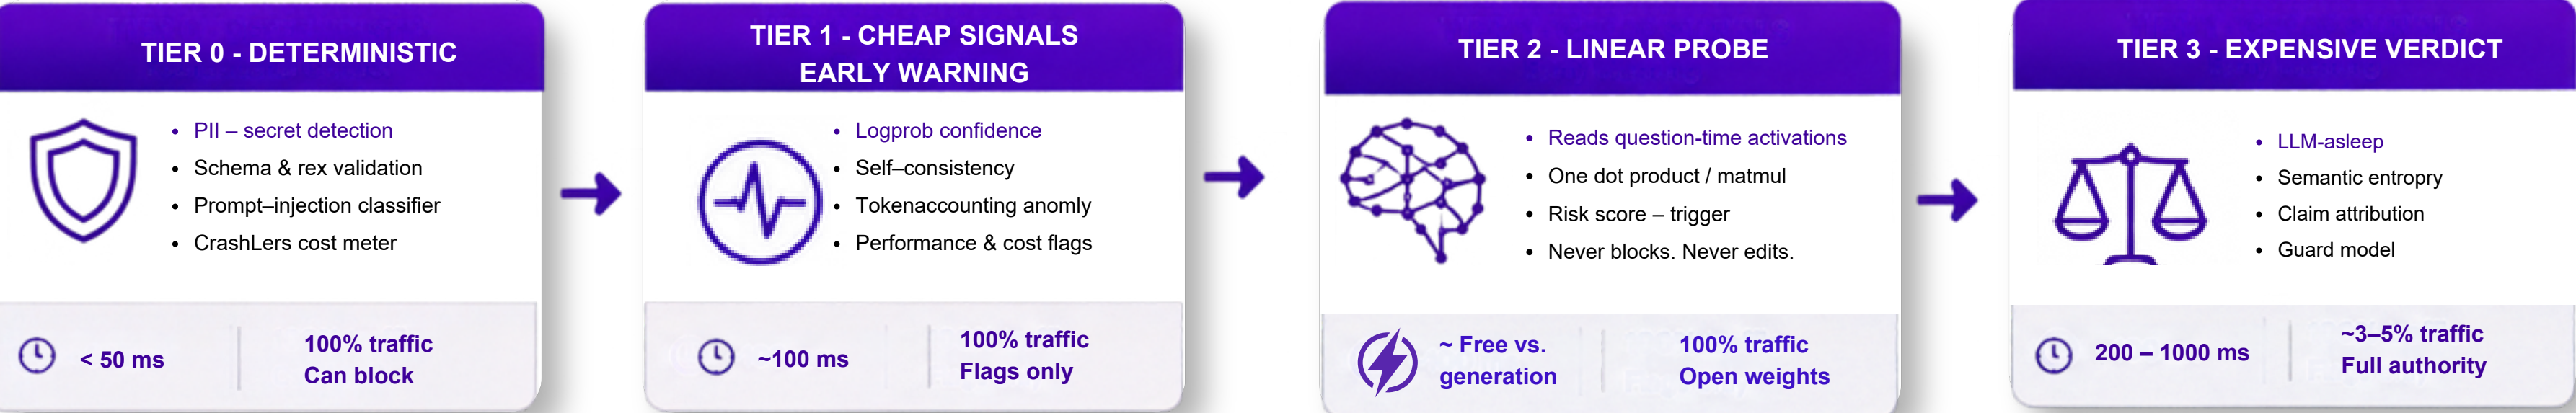


Key challenge:

How do we detect each risk category in real time and decide whether a response should be allowed, edited, blocked, or escalated to a human — without slowing AI down?

> CONTROLPLANE.AI – OUR SOLUTION

Free trigger on 100%. Expensive verdict on -3% – maximize coverage, minimize cost.





https://drive.google.com/file/d/1TwB2H7-A1F57nqFLJE9D7H90q_6nknFa/view?usp=sharing

Thank you

